

# LlaMa

January 12, 2025

```
[1]: # Install necessary libraries
```

```
!pip install transformers yfinance ccxt pandas matplotlib torch
```

Requirement already satisfied: transformers in /usr/local/lib/python3.10/dist-packages (4.47.1)

Requirement already satisfied: yfinance in /usr/local/lib/python3.10/dist-packages (0.2.51)

Collecting ccxt

Downloading ccxt-4.4.47-py2.py3-none-any.whl.metadata (117 kB)

117.7/117.7

kB 5.5 MB/s eta 0:00:00

Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (2.2.2)

Requirement already satisfied: matplotlib in /usr/local/lib/python3.10/dist-packages (3.10.0)

Requirement already satisfied: torch in /usr/local/lib/python3.10/dist-packages (2.5.1+cu121)

Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from transformers) (3.16.1)

Requirement already satisfied: huggingface-hub<1.0,>=0.24.0 in /usr/local/lib/python3.10/dist-packages (from transformers) (0.27.1)

Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.10/dist-packages (from transformers) (1.26.4)

Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from transformers) (24.2)

Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.10/dist-packages (from transformers) (6.0.2)

Requirement already satisfied: regex!=2019.12.17 in /usr/local/lib/python3.10/dist-packages (from transformers) (2024.11.6)

Requirement already satisfied: requests in /usr/local/lib/python3.10/dist-packages (from transformers) (2.32.3)

Requirement already satisfied: tokenizers<0.22,>=0.21 in /usr/local/lib/python3.10/dist-packages (from transformers) (0.21.0)

Requirement already satisfied: safetensors>=0.4.1 in /usr/local/lib/python3.10/dist-packages (from transformers) (0.5.1)

Requirement already satisfied: tqdm>=4.27 in /usr/local/lib/python3.10/dist-packages (from transformers) (4.67.1)

Requirement already satisfied: multitasking>=0.0.7 in /usr/local/lib/python3.10/dist-packages (from yfinance) (0.0.11)

Requirement already satisfied: lxml>=4.9.1 in /usr/local/lib/python3.10/dist-packages (from yfinance) (5.3.0)

Requirement already satisfied: platformdirs>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from yfinance) (4.3.6)

Requirement already satisfied: pytz>=2022.5 in /usr/local/lib/python3.10/dist-packages (from yfinance) (2024.2)

Requirement already satisfied: frozendict>=2.3.4 in /usr/local/lib/python3.10/dist-packages (from yfinance) (2.4.6)

Requirement already satisfied: peewee>=3.16.2 in /usr/local/lib/python3.10/dist-packages (from yfinance) (3.17.8)

Requirement already satisfied: beautifulsoup4>=4.11.1 in /usr/local/lib/python3.10/dist-packages (from yfinance) (4.12.3)

Requirement already satisfied: html5lib>=1.1 in /usr/local/lib/python3.10/dist-packages (from yfinance) (1.1)

Requirement already satisfied: setuptools>=60.9.0 in /usr/local/lib/python3.10/dist-packages (from ccxt) (75.1.0)

Requirement already satisfied: certifi>=2018.1.18 in /usr/local/lib/python3.10/dist-packages (from ccxt) (2024.12.14)

Requirement already satisfied: cryptography>=2.6.1 in /usr/local/lib/python3.10/dist-packages (from ccxt) (43.0.3)

Requirement already satisfied: typing-extensions>=4.4.0 in /usr/local/lib/python3.10/dist-packages (from ccxt) (4.12.2)

Collecting aiohttp<=3.10.11 (from ccxt)

  Downloading aiohttp-3.10.11-cp310-cp310-manylinux\_2\_17\_x86\_64.manylinux2014\_x86\_64.whl.metadata (7.7 kB)

Collecting aiodns>=1.1.1 (from ccxt)

  Downloading aiodns-3.2.0-py3-none-any.whl.metadata (4.0 kB)

Requirement already satisfied: yarl>=1.7.2 in /usr/local/lib/python3.10/dist-packages (from ccxt) (1.18.3)

Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.10/dist-packages (from pandas) (2.8.2)

Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.10/dist-packages (from pandas) (2024.2)

Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (1.3.1)

Requirement already satisfied: cycycler>=0.10 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (0.12.1)

Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (4.55.3)

Requirement already satisfied: kiwisolver>=1.3.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (1.4.8)

Requirement already satisfied: pillow>=8 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (11.1.0)

Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (3.2.1)

Requirement already satisfied: networkx in /usr/local/lib/python3.10/dist-

packages (from torch) (3.4.2)  
 Requirement already satisfied: jinja2 in /usr/local/lib/python3.10/dist-packages (from torch) (3.1.5)  
 Requirement already satisfied: fsspec in /usr/local/lib/python3.10/dist-packages (from torch) (2024.10.0)  
 Requirement already satisfied: sympy==1.13.1 in /usr/local/lib/python3.10/dist-packages (from torch) (1.13.1)  
 Requirement already satisfied: mpmath<1.4,>=1.1.0 in /usr/local/lib/python3.10/dist-packages (from sympy==1.13.1->torch) (1.3.0)  
 Collecting pycares>=4.0.0 (from aiодns>=1.1.1->ccxt)  
   Downloading pycares-4.5.0-cp310-cp310-manylinux\_2\_17\_x86\_64.manylinux2014\_x86\_64.whl.metadata (4.1 kB)  
 Requirement already satisfied: aiohappyeyeballs>=2.3.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp<=3.10.11->ccxt) (2.4.4)  
 Requirement already satisfied: aiosignal>=1.1.2 in /usr/local/lib/python3.10/dist-packages (from aiohttp<=3.10.11->ccxt) (1.3.2)  
 Requirement already satisfied: attrs>=17.3.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp<=3.10.11->ccxt) (24.3.0)  
 Requirement already satisfied: frozenlist>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from aiohttp<=3.10.11->ccxt) (1.5.0)  
 Requirement already satisfied: multidict<7.0,>=4.5 in /usr/local/lib/python3.10/dist-packages (from aiohttp<=3.10.11->ccxt) (6.1.0)  
 Requirement already satisfied: async-timeout<6.0,>=4.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp<=3.10.11->ccxt) (4.0.3)  
 Requirement already satisfied: soupsieve>1.2 in /usr/local/lib/python3.10/dist-packages (from beautifulsoup4>=4.11.1->yfinance) (2.6)  
 Requirement already satisfied: cffi>=1.12 in /usr/local/lib/python3.10/dist-packages (from cryptography>=2.6.1->ccxt) (1.17.1)  
 Requirement already satisfied: six>=1.9 in /usr/local/lib/python3.10/dist-packages (from html5lib>=1.1->yfinance) (1.17.0)  
 Requirement already satisfied: webencodings in /usr/local/lib/python3.10/dist-packages (from html5lib>=1.1->yfinance) (0.5.1)  
 Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests->transformers) (3.4.1)  
 Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests->transformers) (3.10)  
 Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests->transformers) (2.3.0)  
 Requirement already satisfied: propcache>=0.2.0 in /usr/local/lib/python3.10/dist-packages (from yarl>=1.7.2->ccxt) (0.2.1)  
 Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.10/dist-packages (from jinja2->torch) (3.0.2)  
 Requirement already satisfied: pycparser in /usr/local/lib/python3.10/dist-packages (from cffi>=1.12->cryptography>=2.6.1->ccxt) (2.22)  
 Downloading ccxt-4.4.47-py2.py3-none-any.whl (5.6 MB)  
       5.6/5.6 MB  
 81.4 MB/s eta 0:00:00  
 Downloading aiодns-3.2.0-py3-none-any.whl (5.7 kB)

Downloading

aiohttp-3.10.11-cp310-cp310-manylinux\_2\_17\_x86\_64.manylinux2014\_x86\_64.whl (1.2 MB)

1.2/1.2 MB

72.8 MB/s eta 0:00:00

Downloading

pycares-4.5.0-cp310-cp310-manylinux\_2\_17\_x86\_64.manylinux2014\_x86\_64.whl (288 kB)

288.6/288.6 kB

26.3 MB/s eta 0:00:00

Installing collected packages: pycares, aiohttp, aiodns, ccxt

Attempting uninstall: aiohttp

Found existing installation: aiohttp 3.11.11

Uninstalling aiohttp-3.11.11:

Successfully uninstalled aiohttp-3.11.11

Successfully installed aiodns-3.2.0 aiohttp-3.10.11 ccxt-4.4.47 pycares-4.5.0

```
[2]: # Import libraries
import yfinance as yf
import ccxt
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from transformers import LlamaTokenizer, LlamaForCausalLM
import torch

[3]: # Function to fetch stock data using yfinance
def fetch_stock_data(ticker, start_date='2010-01-01', end_date=None):
    stock = yf.Ticker(ticker)
    df = stock.history(start=start_date, end=end_date)
    df = df[['Open', 'High', 'Low', 'Close', 'Volume']].reset_index()
    df.columns = ['timestamp', 'open', 'high', 'low', 'close', 'volume']
    return df

# Function to fetch cryptocurrency data using ccxt
def fetch_crypto_data(symbol='BTC/USD', exchange_name='kraken', timeframe='1d',
    limit=1000):
    exchange_class = getattr(ccxt, exchange_name)
    exchange = exchange_class()
    ohlcv = exchange.fetch_ohlcv(symbol, timeframe=timeframe, limit=limit)
    data = pd.DataFrame(ohlcv, columns=['timestamp', 'open', 'high', 'low',
    'close', 'volume'])
    data['timestamp'] = pd.to_datetime(data['timestamp'], unit='ms')
    return data

# Fetch data for stocks and cryptocurrencies
tesla_data = fetch_stock_data('TSLA', start_date='2020-01-01')
```

```

apple_data = fetch_stock_data('AAPL', start_date='2020-01-01')
btc_data = fetch_crypto_data('BTC/USD')
eth_data = fetch_crypto_data('ETH/USD')

# Display sample data
print("Tesla Data:\n", tesla_data.head())
print("Apple Data:\n", apple_data.head())
print("Bitcoin Data:\n", btc_data.head())
print("Ethereum Data:\n", eth_data.head())

```

Tesla Data:

	timestamp	open	high	low	close \
0	2020-01-02 00:00:00-05:00	28.299999	28.713333	28.114000	28.684000
1	2020-01-03 00:00:00-05:00	29.366667	30.266666	29.128000	29.534000
2	2020-01-06 00:00:00-05:00	29.364668	30.104000	29.333332	30.102667
3	2020-01-07 00:00:00-05:00	30.760000	31.441999	30.224001	31.270666
4	2020-01-08 00:00:00-05:00	31.580000	33.232666	31.215334	32.809334

	volume
0	142981500
1	266677500
2	151995000
3	268231500
4	467164500

Apple Data:

	timestamp	open	high	low	close \
0	2020-01-02 00:00:00-05:00	71.799866	72.856606	71.545380	72.796013
1	2020-01-03 00:00:00-05:00	72.020454	72.851784	71.862915	72.088318
2	2020-01-06 00:00:00-05:00	71.206062	72.701485	70.953995	72.662704
3	2020-01-07 00:00:00-05:00	72.672409	72.929322	72.100418	72.320976
4	2020-01-08 00:00:00-05:00	72.022843	73.787300	72.022843	73.484337

	volume
0	135480400
1	146322800
2	118387200
3	108872000
4	132079200

Bitcoin Data:

	timestamp	open	high	low	close	volume
0	2023-01-24	22926.0	23158.7	22455.9	22633.8	3077.643596
1	2023-01-25	22636.0	23829.3	22320.0	23056.5	5020.204657
2	2023-01-26	23063.2	23293.3	22857.5	23010.6	3753.163921
3	2023-01-27	23010.6	23500.0	22492.8	23077.5	3420.533974
4	2023-01-28	23079.9	23191.9	22900.0	23031.1	1006.279351

Ethereum Data:

	timestamp	open	high	low	close	volume
--	-----------	------	------	-----	-------	--------

0	2023-01-24	1626.40	1640.04	1531.35	1556.01	40939.533459
1	2023-01-25	1555.63	1639.60	1518.00	1611.62	33939.635980
2	2023-01-26	1611.62	1633.39	1578.78	1602.22	20970.688727
3	2023-01-27	1602.14	1621.59	1551.00	1597.31	38471.693576
4	2023-01-28	1597.31	1606.82	1557.24	1572.63	17893.514388

```
[4]: # Function to create a text prompt for LLaMA
def create_text_prompt(data, asset_name="Asset", max_entries=5):
    prompt = f"Here is the historical price data for {asset_name}. Predict the_
    ↪next closing price:\n\n"
    data = data.tail(max_entries)
    for _, row in data.iterrows():
        prompt += f>Date: {row['timestamp'].strftime('%Y-%m-%d')}, Open:_
        ↪{row['open']}, High: {row['high']}, Low: {row['low']}, Close:_
        ↪{row['close']}\n"
    prompt += "\nThe next closing price is: "
    return prompt

# Generate prompts for all four assets
tesla_prompt = create_text_prompt(tesla_data, asset_name="Tesla", max_entries=5)
apple_prompt = create_text_prompt(apple_data, asset_name="Apple", max_entries=5)
btc_prompt = create_text_prompt(btc_data, asset_name="Bitcoin", max_entries=5)
eth_prompt = create_text_prompt(eth_data, asset_name="Ethereum", max_entries=5)

# Display prompts
print("Tesla Prompt:\n", tesla_prompt)
print("\nApple Prompt:\n", apple_prompt)
print("\nBitcoin Prompt:\n", btc_prompt)
print("\nEthereum Prompt:\n", eth_prompt)
```

Tesla Prompt:

Here is the historical price data for Tesla. Predict the next closing price:

Date: 2025-01-03, Open: 381.4800109863281, High: 411.8800048828125, Low: 379.45001220703125, Close: 410.44000244140625  
 Date: 2025-01-06, Open: 423.20001220703125, High: 426.42999267578125, Low: 401.70001220703125, Close: 411.04998779296875  
 Date: 2025-01-07, Open: 405.8299865722656, High: 414.3299865722656, Low: 390.0, Close: 394.3599853515625  
 Date: 2025-01-08, Open: 392.95001220703125, High: 402.5, Low: 387.3999938964844, Close: 394.94000244140625  
 Date: 2025-01-10, Open: 391.3999938964844, High: 399.2799987792969, Low: 377.2900085449219, Close: 394.739990234375

The next closing price is:

Apple Prompt:

Here is the historical price data for Apple. Predict the next closing price:

Date: 2025-01-03, Open: 243.36000061035156, High: 244.17999267578125, Low: 241.88999938964844, Close: 243.36000061035156  
 Date: 2025-01-06, Open: 244.30999755859375, High: 247.3300018310547, Low: 243.1999969482422, Close: 245.0  
 Date: 2025-01-07, Open: 242.97999572753906, High: 245.5500030517578, Low: 241.35000610351562, Close: 242.2100067138672  
 Date: 2025-01-08, Open: 241.9199981689453, High: 243.7100067138672, Low: 240.0500030517578, Close: 242.6999969482422  
 Date: 2025-01-10, Open: 240.00999450683594, High: 240.16000366210938, Low: 233.0, Close: 236.85000610351562

The next closing price is:

Bitcoin Prompt:

Here is the historical price data for Bitcoin. Predict the next closing price:

Date: 2025-01-08, Open: 96932.2, High: 97232.8, Low: 92501.0, Close: 95060.4  
 Date: 2025-01-09, Open: 95060.1, High: 95251.1, Low: 91168.5, Close: 92531.7  
 Date: 2025-01-10, Open: 92523.5, High: 95771.2, Low: 92200.0, Close: 94698.8  
 Date: 2025-01-11, Open: 94698.9, High: 94986.6, Low: 93865.0, Close: 94562.1  
 Date: 2025-01-12, Open: 94562.1, High: 95300.0, Low: 93652.2, Close: 93875.9

The next closing price is:

Ethereum Prompt:

Here is the historical price data for Ethereum. Predict the next closing price:

Date: 2025-01-08, Open: 3380.67, High: 3414.55, Low: 3210.73, Close: 3326.34  
 Date: 2025-01-09, Open: 3326.4, High: 3355.0, Low: 3159.16, Close: 3219.04  
 Date: 2025-01-10, Open: 3218.8, High: 3320.0, Low: 3194.91, Close: 3266.05  
 Date: 2025-01-11, Open: 3266.06, High: 3317.8, Low: 3220.26, Close: 3282.32  
 Date: 2025-01-12, Open: 3281.68, High: 3295.83, Low: 3225.29, Close: 3234.97

The next closing price is:

```
[5]: # Load LLaMA tokenizer and model
model_name = "huggyllama/llama-7b" # Adjust based on available model
tokenizer = LlamaTokenizer.from_pretrained(model_name)
model = LlamaForCausalLM.from_pretrained(model_name, torch_dtype=torch.float16,
↪device_map="auto")

print("LLaMA model loaded successfully.")
```

/usr/local/lib/python3.10/dist-packages/huggingface\_hub/utils/\_auth.py:94:

UserWarning:

The secret `HF\_TOKEN` does not exist in your Colab secrets.

To authenticate with the Hugging Face Hub, create a token in your settings tab

(<https://huggingface.co/settings/tokens>), set it as secret in your Google Colab and restart your session.

You will be able to reuse this secret in all of your notebooks.

Please note that authentication is recommended but still optional to access public models or datasets.

```
warnings.warn(
tokenizer_config.json: 0%|          | 0.00/2.28k [00:00<?, ?B/s]
tokenizer.model: 0%|          | 0.00/500k [00:00<?, ?B/s]
special_tokens_map.json: 0%|          | 0.00/411 [00:00<?, ?B/s]
tokenizer.json: 0%|          | 0.00/1.84M [00:00<?, ?B/s]
```

You are using the default legacy behaviour of the `<class 'transformers.models.llama.tokenization_llama.LlamaTokenizer'>`. This is expected, and simply means that the ``legacy`` (previous) behavior will be used so nothing changes for you. If you want to use the new behaviour, set ``legacy=False``. This should only be set if you understand what it means, and thoroughly read the reason why this was added as explained in <https://github.com/huggingface/transformers/pull/24565> - if you loaded a llama tokenizer from a GGUF file you can ignore this message

```
config.json: 0%|          | 0.00/594 [00:00<?, ?B/s]
model.safetensors.index.json: 0%|          | 0.00/26.8k [00:00<?, ?B/s]
Downloading shards: 0%|          | 0/2 [00:00<?, ?it/s]
model-00001-of-00002.safetensors: 0%|          | 0.00/9.98G [00:00<?, ?B/s]
model-00002-of-00002.safetensors: 0%|          | 0.00/3.50G [00:00<?, ?B/s]
Loading checkpoint shards: 0%|          | 0/2 [00:00<?, ?it/s]
generation_config.json: 0%|          | 0.00/137 [00:00<?, ?B/s]
LLaMA model loaded successfully.
```

```
[6]: # Function to generate predictions using LLaMA
def predict_next_price(prompt):
    inputs = tokenizer(prompt, return_tensors="pt").to('cuda') # Move input to GPU
    with torch.no_grad():
        outputs = model.generate(
            inputs.input_ids,
            max_new_tokens=10,
            pad_token_id=tokenizer.eos_token_id
        )
    prediction = tokenizer.decode(outputs[:, inputs.input_ids.shape[-1]:][0], skip_special_tokens=True)
    return prediction
```



```

# Predict prices for all four assets
predicted_tesla_price = predict_next_price(tesla_prompt)
predicted_apple_price = predict_next_price(apple_prompt)
predicted_btc_price = predict_next_price(btc_prompt)
predicted_eth_price = predict_next_price(eth_prompt)

# Display predictions
print("Predicted Tesla Price:", predicted_tesla_price)
print("Predicted Apple Price:", predicted_apple_price)
print("Predicted Bitcoin Price:", predicted_btc_price)
print("Predicted Ethereum Price:", predicted_eth_price)

```

The attention mask is not set and cannot be inferred from input because pad token is same as eos token. As a consequence, you may observe unexpected behavior. Please pass your input's `attention\_mask` to obtain reliable results.

```

Predicted Tesla Price: 394.739990
Predicted Apple Price: 236.850006
Predicted Bitcoin Price: 93875.9

```

Answer

```
Predicted Ethereum Price: 3234.97
```

Answer

```

[7]: # Function to generate predictions for multiple days
def generate_predictions_for_days(data, num_days=5, asset_name="Asset"):
    future_data = data.copy()
    predictions = []

    for i in range(num_days):
        prompt = create_text_prompt(future_data, asset_name=asset_name,
        ↪max_entries=5)
        predicted_price = predict_next_price(prompt)
        predicted_price_value = float(predicted_price.split()[0])
        predictions.append(predicted_price_value)

    # Add the predicted price to the dataset for the next iteration
    next_date = future_data['timestamp'].iloc[-1] + pd.DateOffset(1)
    new_row = pd.DataFrame({
        'timestamp': [next_date],
        'open': [np.nan],
        'high': [np.nan],
        'low': [np.nan],
        'close': [predicted_price_value],
        'volume': [np.nan]
    })

```

```

        future_data = pd.concat([future_data, new_row], ignore_index=True)

    return predictions

# Generate predictions for all four assets
tesla_predictions = generate_predictions_for_days(tesla_data, num_days=5,
    ↪asset_name="Tesla")
apple_predictions = generate_predictions_for_days(apple_data, num_days=5,
    ↪asset_name="Apple")
btc_predictions = generate_predictions_for_days(btc_data, num_days=5,
    ↪asset_name="Bitcoin")
eth_predictions = generate_predictions_for_days(eth_data, num_days=5,
    ↪asset_name="Ethereum")

# Display predictions
print("Predicted Tesla Prices for Next 5 Days:", tesla_predictions)
print("Predicted Apple Prices for Next 5 Days:", apple_predictions)
print("Predicted Bitcoin Prices for Next 5 Days:", btc_predictions)
print("Predicted Ethereum Prices for Next 5 Days:", eth_predictions)

```

Predicted Tesla Prices for Next 5 Days: [394.73999, 394.73999, 394.73999, 394.73999, 394.73999]

Predicted Apple Prices for Next 5 Days: [236.850006, 236.850006, 236.850006, 236.850006, 236.850006]

Predicted Bitcoin Prices for Next 5 Days: [93875.9, 93875.9, 93875.9, 93875.9, 93875.9]

Predicted Ethereum Prices for Next 5 Days: [3234.97, 3234.97, 3234.97, 3234.97, 3234.97]

```

[8]: # Function to plot predictions along with historical data
def plot_predictions(data, predicted_prices, asset_name="Asset", num_days=5):
    data['timestamp'] = pd.to_datetime(data['timestamp'])
    plt.figure(figsize=(10, 6))

    # Plot historical closing prices
    plt.plot(data['timestamp'], data['close'], label='Historical Prices',
    ↪marker='o', color='blue')

    # Generate future dates
    last_date = data['timestamp'].iloc[-1]
    future_dates = [last_date + pd.DateOffset(days=i+1) for i in
    ↪range(num_days)]


    # Plot predicted prices
    plt.plot(future_dates, predicted_prices, label='Predicted Prices',
    ↪marker='x', color='red', linestyle='--')



```


```
plt.title(f"{asset_name} Historical and Predicted Prices")
plt.xlabel("Date")
plt.ylabel("Price (USD)")
plt.legend()
plt.grid(True)
plt.show()
```

*# Plot predictions for all four assets*

```
plot_predictions(tesla_data, tesla_predictions, asset_name="Tesla", num_days=5)
plot_predictions(apple_data, apple_predictions, asset_name="Apple", num_days=5)
plot_predictions(btc_data, btc_predictions, asset_name="Bitcoin", num_days=5)
plot_predictions(eth_data, eth_predictions, asset_name="Ethereum", num_days=5)
```







```
[9]: # Function to compute evaluation metrics
def compute_metrics(actual, predicted):
    mse = np.mean((np.array(actual) - np.array(predicted)) ** 2)
    rmse = np.sqrt(mse)
    mae = np.mean(np.abs(np.array(actual) - np.array(predicted)))
    return mse, rmse, mae
```

```
[10]: # Predict prices for Tesla
predicted_tesla_price = predict_next_price(tesla_prompt)

# Create a list of identical predictions (as per the original behavior)
predicted_tesla_prices = [float(predicted_tesla_price.split()[0])] * 5

# Compute evaluation metrics for Tesla
actual_tesla_prices = tesla_data['close'].tail(5).tolist()
tesla_mse, tesla_rmse, tesla_mae = compute_metrics(actual_tesla_prices,
    ↪ predicted_tesla_prices)

# Display Tesla metrics
print(f"Tesla - MSE: {tesla_mse}, RMSE: {tesla_rmse}, MAE: {tesla_mae}")
```

```
Tesla - MSE: 102.53816543530188, RMSE: 10.126113046737226, MAE:
6.518005511718764
```

```
[11]: # Predict prices for Apple
predicted_apple_price = predict_next_price(apple_prompt)

# Create a list of identical predictions
predicted_apple_prices = [float(predicted_apple_price.split()[0])] * 5

# Compute evaluation metrics for Apple
actual_apple_prices = apple_data['close'].tail(5).tolist()
apple_mse, apple_rmse, apple_mae = compute_metrics(actual_apple_prices,
    ↪predicted_apple_prices)

# Display Apple metrics
print(f"Apple - MSE: {apple_mse}, RMSE: {apple_rmse}, MAE: {apple_mae}")
```

Apple - MSE: 34.35088675480286, RMSE: 5.860962954566669, MAE: 5.173996075195305

```
[12]: # Predict prices for Bitcoin
predicted_btc_price = predict_next_price(btc_prompt)

# Create a list of identical predictions
predicted_btc_prices = [float(predicted_btc_price.split()[0])] * 5

# Compute evaluation metrics for Bitcoin
actual_btc_prices = btc_data['close'].tail(5).tolist()
btc_mse, btc_rmse, btc_mae = compute_metrics(actual_btc_prices,
    ↪predicted_btc_prices)

# Display Bitcoin metrics
print(f"Bitcoin - MSE: {btc_mse}, RMSE: {btc_rmse}, MAE: {btc_mae}")
```

Bitcoin - MSE: 871589.7480000046, RMSE: 933.5897107402184, MAE: 807.56000000000035

```
[13]: # Predict prices for Ethereum
predicted_eth_price = predict_next_price(eth_prompt)

# Create a list of identical predictions
predicted_eth_prices = [float(predicted_eth_price.split()[0])] * 5

# Compute evaluation metrics for Ethereum
actual_eth_prices = eth_data['close'].tail(5).tolist()
eth_mse, eth_rmse, eth_mae = compute_metrics(actual_eth_prices,
    ↪predicted_eth_prices)

# Display Ethereum metrics
print(f"Ethereum - MSE: {eth_mse}, RMSE: {eth_rmse}, MAE: {eth_mae}")
```

Ethereum - MSE: 2362.046140000023, RMSE: 48.60088620591216, MAE:

37.146000000000186

[ ]: